

Chapter 19 – Parallel Processing Part IV

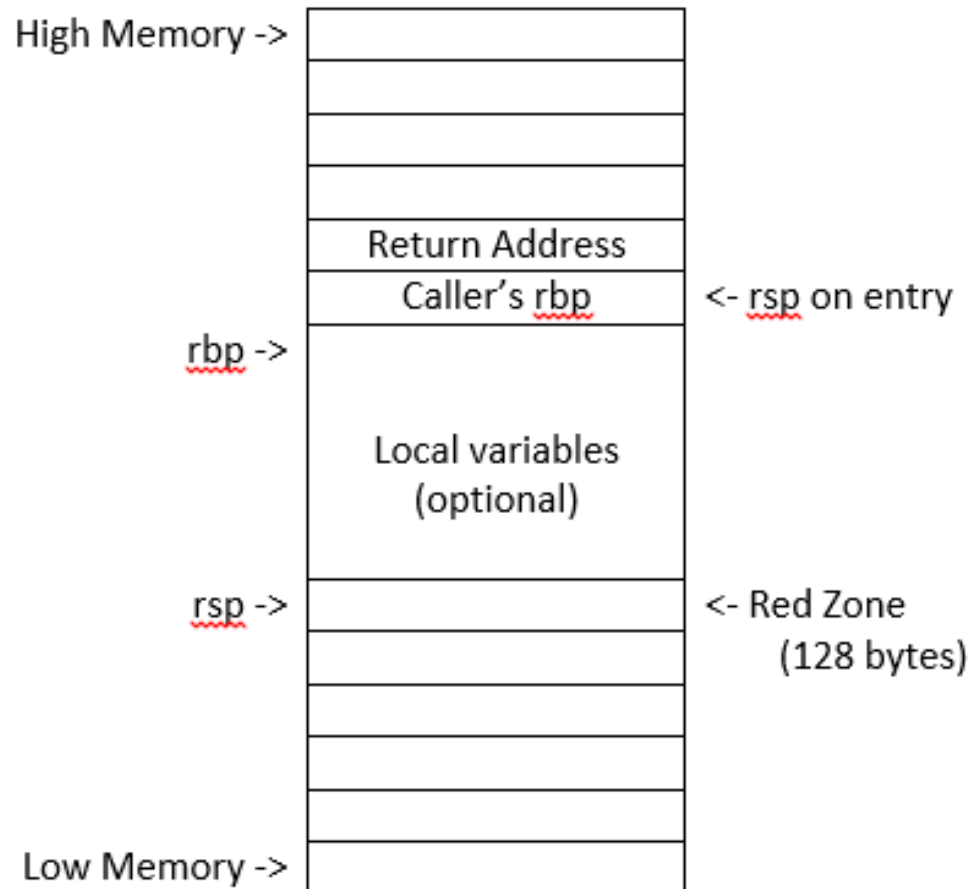


CPSC 240-09
John Overton

Review

More Parallel Processing...

- The Red Zone: An area of 128 bytes that is before the `rsp` pointer (remember, the `rsp` grows down)



```
_start:
    push    rbp
    mov     rbp, rsp
    sub     rsp, 0x10 ; allocate local space
    ...
```

```
; Allocating local variables: 2 qword sized variables
var1     equ     -0x8
var2     equ     -0x10

; using local variables:
    mov     qword [rbp+var1], rax ; save rax
    mov     rbx, qword [rbp+var2]
```

Parallel Processing – Atomic instructions in x86_64

Atomicity in x86_64

- The lock prefix can be used in front of several different instructions to cause the instruction to be an atomic operation. Can be used in front of:
 - ADD, SUB, INC, DEC, NEG, NOT, OR, AND, CMPXCHG, 10 others
- We could use lock in front of an add instruction in our Assignment #9 where each thread added its nbrFactors to a grandTotal variable (instead of having an array) –

```
-- r11 is nbrFactors  
;mov [thread_total+(r8*8)], r11  
lock add qword [grand_total], r11
```

- This is where I demo some code that uses lock in front of an inc instruction

Compare-and-Swap Method of Atomicity

Using the `cmpxchg` instruction with the lock prefix is a true compare-and-swap atomic operation

And example of `cmpxchg`:

```
lock cmpxchg  qword [lockword], r8
```

- Compares the contents of `rax` with the contents of `lockword`
 - If equal, then the contents of `r8` is copied to `lockword`.
 - If not equal, then the contents of `lockword` are copied to `rax` to get ready to retry the `cmpxchg` instruction
- The `lock` prefix makes the instruction atomic

Compare-and-Swap Method of Atomicity – Example: spinlock

```
mutex    dq    0
...
mov      rdi, mutex
call     spinlock
...

; -----
; Lock a mutex.  On entry, rdi will point to qword lock word.
; This routine will not return until it "owns" the mutex.
; The mutex will only contain 0 (unlocked) or 1 (locked)
spinlock:
    mov     rax, 0          ; what we expect in the mutex
    mov     r8, 1           ; what we want in the mutex
lockloop:
    lock    cmpxchg qword [rdi], r8 ; lock prefix to cmpxchg
    ; if equal, r8->[rdi]
    jne     lockloop        ; if not equal, [rdi]->rax
    ret

; -----
; Unlock a mutex.  Assume we were the thread that locked it.
; On entry, rdi will point to qword lock word.
spinunlk:
    mov     r8, 0
    xchg    r8, qword [rdi] ; no need to use lock - implied.
    ret
```

Phase 3 - Results running with NUMBER_TO_FACTOR = 2 Billion

Total factors of 2000000000 = 110

With 1 thread:

real 0m14.700s
user 0m14.656s

real = Elapsed time
User = time processor(s) spend working

With 4 threads:

real 0m4.373s
user 0m17.188s ← Increase in overhead

With 8 threads:

real 0m2.745s ← dramatic decrease in elapsed time
user 0m20.234s

With 16 threads:

real 0m2.708s ← no decrease in elapsed time because I only have 8 processors
user 0m20.234s

Assignment – Phase 3 – with “lock add”

In Phase 3, you will replace your Hello World part of your threads with an equation to factor a number. That is, given a number, count the number of numbers that can be wholly divided into that number:

```
NUMBER_TO_FACTOR    equ    10000000000
RANGE_EACH_THREAD    equ    NUMBER_TO_FACTOR / MAX_THREADS
```

```
from = (threadIndex * RANGE_EACH_THREAD) + 1
to = from + RANGE_EACH_THREAD
```

```
for( i = from; i < to; i++ ) {
    if ( number % i == 0 )
        nbrFactors++
}
```

Assignment – Phase 3 – with “lock add”

In your thread function, it's easier to use registers for all of the variables since each thread gets its own registers.

On Tuesday, we talked about a solution that involved each thread saving `nbrFactors` in its own element in an array.

```
threadTotals[threadIndex] = nbrFactors
```

We can now use the “lock add” solution to add `nbrFactors` to a grand total without worrying about each thread causing a race condition with the grand total variable